

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    classification_report, confusion_matrix, roc_auc_score,
    roc_curve, precision_recall_curve, average_precision_score
)
import pickle
import os
import warnings
warnings.filterwarnings('ignore')

# =====
# CONFIGURACIÓN
# =====
TEST_SIZE = 0.2
RANDOM_STATE = 42
CV_FOLDS = 5
N_ESTIMATORS = 100
MAX_DEPTH = 15
MIN_SAMPLES_SPLIT = 10
MIN_SAMPLES_LEAF = 5

# Directorio de salida
OUTPUT_DIR = "../../data/random-forest-output"
os.makedirs(OUTPUT_DIR, exist_ok=True)

print("=*80")
print("MODELO DE RANDOM FOREST - HOME CREDIT DEFAULT RISK")
print("=*80")
print(f"\nTipo de problema: CLASIFICACIÓN BINARIA (0 = Pago, 1 = Default)")
print(f"Criterio de split: GINI (mide impureza para clasificación)")
print(f"Directorio de salida: {OUTPUT_DIR}")

=====
```

```
=====
=====
MODELO DE RANDOM FOREST - HOME CREDIT DEFAULT RISK
=====
=====
```

```
Tipo de problema: CLASIFICACIÓN BINARIA (0 = Pago, 1 = Default)
Criterio de split: GINI (mide impureza para clasificación)
Directorio de salida: ../../data/random-forest-output
```

```
In [2]: # =====
# 1. CARGA Y VERIFICACIÓN DE DATOS
# =====
print("\n[1/15] Cargando y verificando datos...")

df = pd.read_csv("../../data/processed/variables.csv")
```

```
print(f"Dataset cargado: {df.shape}")
print(f"Usuarios: {len(df)}")
print(f"Variables: {len(df.columns)}")
print(f"Tasa de default: {df['TARGET'].mean()*100:.2f}%")
```

```
[1/15] Cargando y verificando datos...
Dataset cargado: (307511, 41)
Usuarios: 307,511
Variables: 41
Tasa de default: 8.07%
```

```
In [3]: # =====
# 2. PREPARACIÓN DE DATOS
# =====
print("\n[2/15] Preparando datos...")

X = df.drop(['SK_ID_CURR', 'TARGET'], axis=1)
y = df['TARGET']

print(f"Variables predictoras: {X.shape[1]}")
print(f"Variable objetivo: {y.name}")

numeric_vars = X.select_dtypes(include=[np.number]).columns.tolist()
categorical_vars = X.select_dtypes(include=['object']).columns.tolist()

print(f"Variables numericas: {len(numeric_vars)}")
print(f"Variables categoricas: {len(categorical_vars)}")
```

[2/15] Preparando datos...
Variables predictoras: 39
Variable objetivo: TARGET
Variables numericas: 34
Variables categoricas: 5

```
In [ ]: # =====
# 3. MANEJO DE VALORES FALTANTES
# =====
print("\n[3/15] Analizando valores faltantes...")

missing_pct = (X.isnull().sum() / len(X) * 100).sort_values(ascending=False)
missing_vars = missing_pct[missing_pct > 0]

if len(missing_vars) > 0:
    print(f"Variables con valores faltantes: {len(missing_vars)}")
    print(missing_vars.head(10))

    print("Imputando valores faltantes...")
    for col in numeric_vars:
        if X[col].isnull().sum() > 0:
            X[col].fillna(X[col].median(), inplace=True)

    for col in categorical_vars:
        if X[col].isnull().sum() > 0:
            X[col].fillna(X[col].mode()[0], inplace=True)

    print("Valores faltantes imputados")
```

```
    else:  
        print("No hay valores faltantes")
```

```
[3/15] Analizando valores faltantes...  
Variables con valores faltantes: 10  
TASA_INTERES_PROMEDIO      98.501192  
RATIO_PAGO_MINIMO_TC       80.726218  
EXT_SOURCE_1                 56.381073  
EXT_SOURCE_3                 19.825307  
PLAZO_PROMEDIO                5.485657  
RATIO_PAGO_CUOTA              5.163718  
EXT_SOURCE_2                 0.214626  
SCORE_PROMEDIO                0.055933  
AMT_ANNUITY                  0.003902  
CNT_FAM_MEMBERS               0.000650  
dtype: float64  
Imputando valores faltantes...  
Valores faltantes imputados
```

```
In [5]: # ======  
# 4. CODIFICACIÓN DE VARIABLES CATEGÓRICAS  
# ======  
print("\n[4/15] Codificando variables categoricas...")  
  
label_encoders = {}  
  
if len(categorical_vars) > 0:  
    for col in categorical_vars:  
        le = LabelEncoder()  
        X[col] = le.fit_transform(X[col].astype(str))  
        label_encoders[col] = le  
  
    print(f"{len(categorical_vars)} variables codificadas")  
    print(f"Clases: {categorical_vars}")
```

```
[4/15] Codificando variables categoricas...  
5 variables codificadas  
Clases: ['NAME_FAMILY_STATUS', 'CODE_GENDER', 'NAME_EDUCATION_TYPE', 'FLAG_OWN_CAR', 'FLAG_OWN_REALTY']
```

```
In [6]: # =====
# 5. DIVISIÓN TRAIN/TEST
# =====
print("\n[5/15] Dividiendo datos en Train/Test...")

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=TEST_SIZE,
    random_state=RANDOM_STATE,
    stratify=y
)

print(f"Train set: {X_train.shape[0]} usuarios ({X_train.shape[0]/len(X)*100}%)")
print(f"Test set: {X_test.shape[0]} usuarios ({X_test.shape[0]/len(X)*100}%)")
print(f"\nDistribucion TARGET en Train:")
print(f"  Clase 0 (Pago): {(y_train==0).sum()} ({(y_train==0).sum()}/len(y_train))")
print(f"  Clase 1 (Default): {(y_train==1).sum()} ({(y_train==1).sum()}/len(y_train))")

[5/15] Dividiendo datos en Train/Test...
Train set: 246,008 usuarios (80.0%)
Test set: 61,503 usuarios (20.0%)

Distribucion TARGET en Train:
  Clase 0 (Pago): 226,148 (91.93%)
  Clase 1 (Default): 19,860 (8.07%)
```

```
In [7]: # =====
# 6. NOTA SOBRE NORMALIZACIÓN
# =====
print("\n[6/15] Nota sobre normalizacion...")

# Random Forest NO requiere normalización ya que es un modelo basado en árboles de decisión
# Los árboles de decisión son invariantes a transformaciones monótonas de las variables
# Usamos los datos sin escalar para mejor interpretabilidad

print("Random Forest no requiere normalizacion de datos")
print("Los arboles de decision son invariantes a escalas de las variables")
print("Se usaran los datos originales (sin escalar)")

[6/15] Nota sobre normalizacion...
Random Forest no requiere normalizacion de datos
Los arboles de decision son invariantes a escalas de las variables
Se usaran los datos originales (sin escalar)
```

Random Forest 1

```
In [8]: # =====
# 7. ENTRENAMIENTO DEL MODELO
# =====
print("\n[7/15] Entrenando Random Forest...")

model = RandomForestClassifier(
    n_estimators=N_ESTIMATORS,
    max_depth=MAX_DEPTH,
```

```

        min_samples_split=MIN_SAMPLES_SPLIT,
        min_samples_leaf=MIN_SAMPLES_LEAF,
        class_weight='balanced',
        random_state=RANDOM_STATE,
        n_jobs=-1,
        verbose=1
    )

model.fit(X_train, y_train)

print("\nModelo entrenado exitosamente")
print(f"Numero de arboles: {model.n_estimators}")
print(f"Profundidad maxima: {model.max_depth}")
print(f"Min samples split: {model.min_samples_split}")
print(f"Min samples leaf: {model.min_samples_leaf}")
print(f"Clases: {model.classes_}")

```

[7/15] Entrenando Random Forest...

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 4.7s

Modelo entrenado exitosamente

Numero de arboles: 100

Profundidad maxima: 15

Min samples split: 10

Min samples leaf: 5

Clases: [0 1]

[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 12.4s finished

In [9]:

```

# =====
# 8. VALIDACIÓN CRUZADA
# =====

print(f"\n[8/15] Validacion Cruzada ({CV_FOLDS}-fold en Train)...")

cv_scores = cross_val_score(
    model, X_train, y_train,
    cv=CV_FOLDS,
    scoring='roc_auc',
    n_jobs=-1
)

print(f"ROC-AUC por fold: {cv_scores}")
print(f"Media: {cv_scores.mean():.4f} (+/- {cv_scores.std():.4f})")

```

[8/15] Validacion Cruzada (5-fold en Train)...

```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:  15.7s
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:  15.7s
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:  16.0s
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:  16.0s
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:  16.0s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  40.5s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done  34 tasks      | elapsed:  0.0s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  40.7s finished
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  40.7s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done  34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  40.8s finished
[Parallel(n_jobs=8)]: Done  34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done  34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.2s finished
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  41.0s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.3s finished
[Parallel(n_jobs=8)]: Done  34 tasks      | elapsed:  0.0s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.2s finished
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
ROC-AUC por fold: [0.74124295 0.7320359  0.74449844 0.7443346  0.74670445]
Media: 0.7418 (+/- 0.0052)

```

In [10]:

```

# =====
# 9. PREDICCIONES
# =====
print("\n[9/15] Generando predicciones...")

y_train_pred = model.predict(X_train)
y_train_proba = model.predict_proba(X_train)[:, 1]

y_test_pred = model.predict(X_test)
y_test_proba = model.predict_proba(X_test)[:, 1]

print("Predicciones generadas")

```

[9/15] Generando predicciones...

```
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.  
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.2s  
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.4s finished  
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.  
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.2s  
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.4s finished  
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.  
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s  
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished  
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.  
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s  
Predicciones generadas  
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
```

```
In [11]: # =====  
# 10. EVALUACIÓN EN TRAIN  
# =====  
print("\n[10/15] Evaluando modelo en TRAIN SET...")  
print("=*80)  
  
train_auc = roc_auc_score(y_train, y_train_proba)  
train_ap = average_precision_score(y_train, y_train_proba)  
  
print(f"\nMetricas Globales:")  
print(f"ROC-AUC Score: {train_auc:.4f}")  
print(f"Average Precision: {train_ap:.4f}")  
  
print(f"\nClassification Report:")  
print(classification_report(y_train, y_train_pred,  
                           target_names=['Pago (0)', 'Default (1)']))  
  
print(f"\nConfusion Matrix:")  
cm_train = confusion_matrix(y_train, y_train_pred)  
print(cm_train)  
print(f"\nInterpretacion:")  
print(f" Verdaderos Negativos (TN): {cm_train[0,0]} - Predijo 'Pago' y si  
print(f" Falsos Positivos (FP): {cm_train[0,1]} - Predijo 'Default' pero  
print(f" Falsos Negativos (FN): {cm_train[1,0]} - Predijo 'Pago' pero hizo  
print(f" Verdaderos Positivos (TP): {cm_train[1,1]} - Predijo 'Default' y
```

```
[10/15] Evaluando modelo en TRAIN SET...
```

```
=====
```

```
=====
```

Metricas Globales:

ROC-AUC Score: 0.9222

Average Precision: 0.5242

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Pago (0) | 0.98 | 0.87 | 0.92 | 226148 |
| Default (1) | 0.36 | 0.78 | 0.49 | 19860 |
| accuracy | | | 0.87 | 246008 |
| macro avg | 0.67 | 0.83 | 0.71 | 246008 |
| weighted avg | 0.93 | 0.87 | 0.89 | 246008 |

Confusion Matrix:

```
[[197868 28280]
```

```
 [ 4282 15578]]
```

Interpretacion:

Verdaderos Negativos (TN): 197,868 – Predijo 'Pago' y si pago

Falsos Positivos (FP): 28,280 – Predijo 'Default' pero pago

Falsos Negativos (FN): 4,282 – Predijo 'Pago' pero hizo default [CRITICO]

Verdaderos Positivos (TP): 15,578 – Predijo 'Default' y si hizo default

```
In [12]:
```

```
# =====
# 11. EVALUACIÓN EN TEST
# =====

print("\n[11/15] Evaluando modelo en TEST SET (metricas definitivas)...")
print("=*80)

test_auc = roc_auc_score(y_test, y_test_proba)
test_ap = average_precision_score(y_test, y_test_proba)

print(f"\nMetricas Globales:")
print(f"ROC-AUC Score: {test_auc:.4f}")
print(f"Average Precision: {test_ap:.4f}")

print(f"\nClassification Report:")
print(classification_report(y_test, y_test_pred,
                           target_names=['Pago (0)', 'Default (1)']))

print(f"\nConfusion Matrix:")
cm_test = confusion_matrix(y_test, y_test_pred)
print(cm_test)
print(f"\nInterpretacion:")
print(f" Verdaderos Negativos (TN): {cm_test[0,0]} - Predijo 'Pago' y si
print(f" Falsos Positivos (FP): {cm_test[0,1]} - Predijo 'Default' pero p
print(f" Falsos Negativos (FN): {cm_test[1,0]} - Predijo 'Pago' pero hizo
print(f" Verdaderos Positivos (TP): {cm_test[1,1]} - Predijo 'Default' y
```

```
[11/15] Evaluando modelo en TEST SET (metricas definitivas)...
```

```
=====
```

Metricas Globales:

ROC-AUC Score: 0.7451

Average Precision: 0.2207

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Pago (0) | 0.95 | 0.86 | 0.90 | 56538 |
| Default (1) | 0.21 | 0.44 | 0.29 | 4965 |
| accuracy | | | 0.82 | 61503 |
| macro avg | 0.58 | 0.65 | 0.59 | 61503 |
| weighted avg | 0.89 | 0.82 | 0.85 | 61503 |

Confusion Matrix:

```
[[48497  8041]
 [ 2777 2188]]
```

Interpretacion:

Verdaderos Negativos (TN): 48,497 – Predijo 'Pago' y si pago

Falsos Positivos (FP): 8,041 – Predijo 'Default' pero pago

Falsos Negativos (FN): 2,777 – Predijo 'Pago' pero hizo default [CRITICO]

Verdaderos Positivos (TP): 2,188 – Predijo 'Default' y si hizo default

In [13]:

```
# =====
# 12. COMPARACIÓN TRAIN VS TEST
# =====
print("\n[12/15] Comparando TRAIN vs TEST (analisis de overfitting)...")
print("=*80)

print(f"\n{'Metrica':<25} {'Train':<12} {'Test':<12} {'Diferencia':<12}")
print("-" * 65)
print(f"{'ROC-AUC':<25} {train_auc:<12.4f} {test_auc:<12.4f} {abs(train_auc-test_auc):<12.4f}")
print(f"{'Average Precision':<25} {train_ap:<12.4f} {test_ap:<12.4f} {abs(train_ap-test_ap):<12.4f}")

diff_auc = abs(train_auc - test_auc)
if diff_auc < 0.02:
    print("\nExcelente generalizacion (diferencia < 2%)")
elif diff_auc < 0.05:
    print("\nBuena generalizacion (diferencia < 5%)")
elif diff_auc < 0.10:
    print("\nPosible ligero overfitting (diferencia 5-10%)")
else:
    print("\nOverfitting detectado (diferencia > 10%)")
```

[12/15] Comparando TRAIN vs TEST (analisis de overfitting)...

=====

| Metrica | Train | Test | Diferencia |
|-------------------|--------|--------|------------|
| ROC-AUC | 0.9222 | 0.7451 | 0.1771 |
| Average Precision | 0.5242 | 0.2207 | 0.3035 |

Overfitting detectado (diferencia > 10%)

In [14]:

```
# =====
# 13. IMPORTANCIA DE VARIABLES
# =====
print("\n[13/15] Calculando importancia de variables...")
print("=*80)

# Random Forest usa importancia basada en Gini (mean decrease in impurity)
# GINI se usa porque es un problema de CLASIFICACIÓN (predecir 0 o 1)
# - Gini mide la impureza de los nodos: que tan mezcladas están las clases
# - Un nodo "puro" tiene solo una clase (Gini = 0)
# - Un nodo con 50/50 tiene máxima impureza (Gini = 0.5 para binario)

feature_importance = pd.DataFrame({
    'Variable': X.columns,
    'Importancia': model.feature_importances_
}).sort_values('Importancia', ascending=False)

print("\nTop 20 variables mas importantes para predecir default:\n")
print(feature_importance.head(20).to_string(index=False))

print("\nInterpretacion de importancia (Gini Importance):")
print(" - Problema de CLASIFICACIÓN BINARIA: predecir 0 (Pago) o 1 (Default")
print(" - Gini mide cuanto reduce cada variable la impureza en los splits")
print(" - Valores mas altos = mayor capacidad para separar las clases")
print(" - Basado en reduccion promedio de impureza en los 100 arboles")
```

[13/15] Calculando importancia de variables...

```
=====
```

Top 20 variables mas importantes para predecir default:

| Variable | Importancia |
|-------------------------|-------------|
| SCORE_PROMEDIO | 0.183249 |
| EXT_SOURCE_2 | 0.095282 |
| EXT_SOURCE_3 | 0.094086 |
| DAYS_BIRTH | 0.042976 |
| EDAD_ANOS | 0.042558 |
| EXT_SOURCE_1 | 0.037583 |
| RATIO_PAGO_CUOTA | 0.036051 |
| AMT_ANNUITY | 0.035872 |
| AMT_CREDIT | 0.034971 |
| CREDIT_INCOME_RATIO | 0.033973 |
| PLAZO_PROMEDIO | 0.031336 |
| TOTAL_CREDITO_OTORGADO | 0.030732 |
| TOTAL_CREDITO_HISTORICO | 0.030321 |
| MONTO_PROMEDIO_PREVIO | 0.030230 |
| TOTAL_DEUDA_ACTUAL | 0.024201 |
| INGRESO_PER_CAPITA | 0.023310 |
| AMT_INCOME_TOTAL | 0.021557 |
| CREDITOS_CERRADOS | 0.016803 |
| RATIO_PAGO_MINIMO_TC | 0.016604 |
| NUM_PRESTAMOS_PREVIOS | 0.016398 |

Interpretacion de importancia (Gini Importance):

- Problema de CLASIFICACIÓN BINARIA: predecir 0 (Pago) o 1 (Default)
- Gini mide cuanto reduce cada variable la impureza en los splits
- Valores mas altos = mayor capacidad para separar las clases
- Basado en reduccion promedio de impureza en los 100 arboles

In [15]:

```
# =====
# 14. VISUALIZACIONES
# =====
print("\n[14/15] Generando visualizaciones...")

# Calcular deciles
test_results = pd.DataFrame({
    'y_true': y_test,
    'y_proba': y_test_proba
})
test_results['decil'] = pd.qcut(test_results['y_proba'], q=10, labels=False)

# Calcular TPR por decil
decil_stats = test_results.groupby('decil').agg({
    'y_true': ['sum', 'count', 'mean']
}).reset_index()
decil_stats.columns = ['decil', 'positivos', 'total', 'tpr']
decil_stats = decil_stats.sort_values('decil', ascending=False)

# Crear gráficas
fig, axes = plt.subplots(2, 2, figsize=(16, 14))
```

```

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_test_proba)
axes[0, 0].plot(fpr, tpr, linewidth=2, color='forestgreen', label=f'ROC (AUC')
axes[0, 0].plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random')
axes[0, 0].set_xlabel('False Positive Rate', fontsize=12)
axes[0, 0].set_ylabel('True Positive Rate', fontsize=12)
axes[0, 0].set_title('ROC Curve – Test Set (Random Forest)', fontsize=14, fontweight='bold')
axes[0, 0].legend()
axes[0, 0].grid(alpha=0.3)

# Precision-Recall Curve
precision, recall, thresholds_pr = precision_recall_curve(y_test, y_test_proba)
sort_idx = np.argsort(recall)
recall_sorted = recall[sort_idx]
precision_sorted = precision[sort_idx]

axes[0, 1].plot(recall_sorted, precision_sorted, linewidth=2, color='forestgreen')
axes[0, 1].axhline(y=y_test.mean(), color='gray', linestyle='--', linewidth=1)
axes[0, 1].set_xlabel('Recall (TPR)', fontsize=12)
axes[0, 1].set_ylabel('Precision', fontsize=12)
axes[0, 1].set_title('Precision-Recall Curve – Test Set', fontsize=14, fontweight='bold')
axes[0, 1].set_xlim([0, 1])
axes[0, 1].set_ylim([0, 1])
axes[0, 1].legend()
axes[0, 1].grid(alpha=0.3)

# TPR por Decil
axes[1, 0].bar(decil_stats['decil'], decil_stats['tpr'], color='forestgreen')
axes[1, 0].set_xlabel('Decil de Score (10 = Mayor Riesgo)', fontsize=12)
axes[1, 0].set_ylabel('True Positive Rate (TPR)', fontsize=12)
axes[1, 0].set_title('TPR por Decil de Probabilidad', fontsize=14, fontweight='bold')
axes[1, 0].set_xticks(range(1, 11))
axes[1, 0].grid(alpha=0.3, axis='y')
axes[1, 0].set_ylim(0, 1)

for i, row in decil_stats.iterrows():
    axes[1, 0].text(row['decil'], row['tpr'] + 0.02, f'{row["tpr"]:.2f}', ha='center', va='bottom', fontsize=10)

# Feature Importance (Top 15)
top_features = feature_importance.head(15)
axes[1, 1].barh(range(len(top_features)), top_features['Importancia'].values)
axes[1, 1].set_yticks(range(len(top_features)))
axes[1, 1].set_yticklabels(top_features['Variable'].values)
axes[1, 1].set_xlabel('Importancia (Gini)', fontsize=12)
axes[1, 1].set_title('Top 15 Variables Mas Importantes', fontsize=14, fontweight='bold')
axes[1, 1].invert_yaxis()
axes[1, 1].grid(alpha=0.3, axis='x')

plt.tight_layout()
plt.savefig(f'{OUTPUT_DIR}/random_forest_curves.png', dpi=300, bbox_inches='tight')
print(f"Graficas guardadas: {OUTPUT_DIR}/random_forest_curves.png")

plt.show()

# Imprimir tabla de deciles

```

```

print("\nTabla de TPR por Decil:")
print(decil_stats.to_string(index=False))

```

[14/15] Generando visualizaciones...

Graficas guardadas: ../../data/random-forest-output/random_forest_curves.png

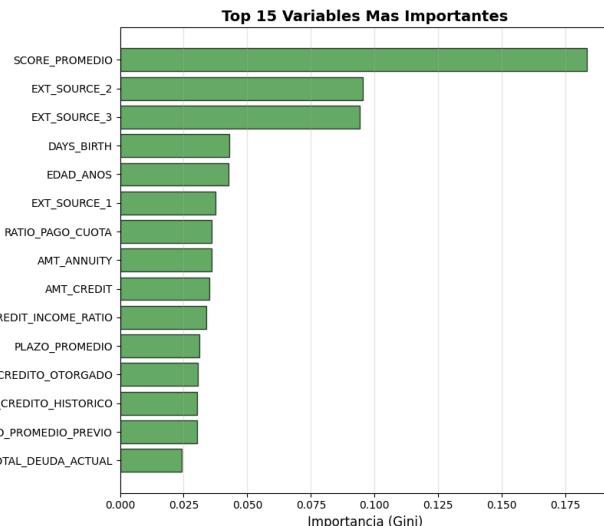
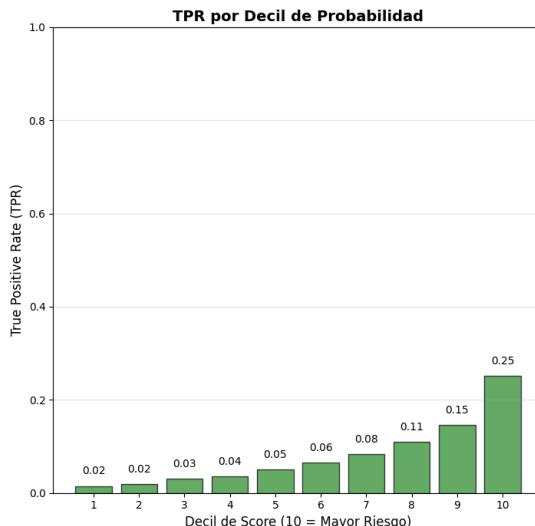
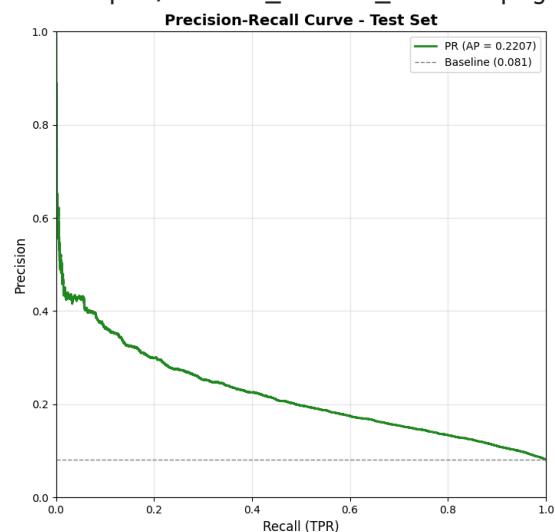
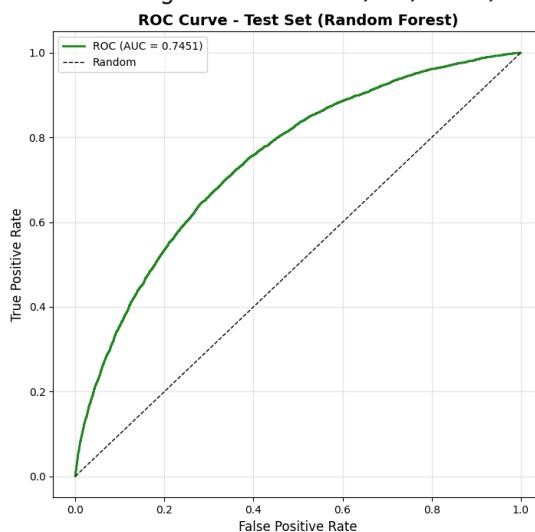


Tabla de TPR por Decil:

| decil | positivos | total | tpr |
|-------|-----------|-------|----------|
| 10 | 1546 | 6151 | 0.251341 |
| 9 | 899 | 6150 | 0.146179 |
| 8 | 679 | 6150 | 0.110407 |
| 7 | 510 | 6150 | 0.082927 |
| 6 | 399 | 6150 | 0.064878 |
| 5 | 309 | 6151 | 0.050236 |
| 4 | 219 | 6150 | 0.035610 |
| 3 | 191 | 6150 | 0.031057 |
| 2 | 120 | 6150 | 0.019512 |
| 1 | 93 | 6151 | 0.015119 |

In [16]:

```

# =====
# 15. GUARDAR MODELO Y RESULTADOS
# =====
print("\n[15/15] Guardando modelo y resultados...")

with open(f'{OUTPUT_DIR}/random_forest_model.pkl', 'wb') as f:
    pickle.dump(model, f)

```

```

with open(f'{OUTPUT_DIR}/label_encoders_rf.pkl', 'wb') as f:
    pickle.dump(label_encoders, f)

feature_importance.to_csv(f'{OUTPUT_DIR}/feature_importance_rf.csv', index=False)

test_predictions = pd.DataFrame({
    'SK_ID_CURR': df.iloc[X_test.index]['SK_ID_CURR'].values,
    'TARGET_Real': y_test.values,
    'TARGET_Predicho': y_test_pred,
    'Probabilidad_Default': y_test_proba
})
test_predictions.to_csv(f'{OUTPUT_DIR}/test_predictions_rf.csv', index=False)

print(f"Archivos guardados en: {OUTPUT_DIR}/")
print(" - random_forest_model.pkl")
print(" - label_encoders_rf.pkl")
print(" - feature_importance_rf.csv")
print(" - test_predictions_rf.csv")
print(" - random_forest_curves.png")

# =====
# RESUMEN FINAL
# =====

print("\n" + "="*80)
print("ANALISIS COMPLETADO – RANDOM FOREST")
print("="*80)
print(f"\nTipo de problema: CLASIFICACIÓN BINARIA")
print(f" - Objetivo: Predecir si un cliente hará default (1) o pagará (0)")
print(f" - Criterio de split: GINI (mide impureza para clasificación)")
print(f"\nResultados Finales (TEST SET):")
print(f" ROC-AUC: {test_auc:.4f}")
print(f" Average Precision: {test_ap:.4f}")
print(f" Accuracy: {(y_test_pred == y_test).mean():.4f}")
print(f"\nConfiguracion del modelo:")
print(f" N_ESTIMATORS: {N_ESTIMATORS}")
print(f" MAX_DEPTH: {MAX_DEPTH}")
print(f" MIN_SAMPLES_SPLIT: {MIN_SAMPLES_SPLIT}")
print(f" MIN_SAMPLES_LEAF: {MIN_SAMPLES_LEAF}")
print(f"\nArchivos guardados en: {OUTPUT_DIR}/")
print(f"\nEl modelo esta listo para produccion")
print("="*80)

```

```
[15/15] Guardando modelo y resultados...
Archivos guardados en: ../../data/random-forest-output/
- random_forest_model.pkl
- label_encoders_rf.pkl
- feature_importance_rf.csv
- test_predictions_rf.csv
- random_forest_curves.png

=====
=====
ANALISIS COMPLETADO - RANDOM FOREST
=====
=====

Tipo de problema: CLASIFICACIÓN BINARIA
- Objetivo: Predecir si un cliente hará default (1) o pagará (0)
- Criterio de split: GINI (mide impureza para clasificación)
```

Resultados Finales (TEST SET):
ROC-AUC: 0.7451
Average Precision: 0.2207
Accuracy: 0.8241

Configuracion del modelo:
N_ESTIMATORS: 100
MAX_DEPTH: 15
MIN_SAMPLES_SPLIT: 10
MIN_SAMPLES_LEAF: 5

Archivos guardados en: ../../data/random-forest-output/

El modelo esta listo para produccion

=====

Random forest 2

```
In [17]: # Drop de variables redundantes/derivadas
df = df.drop(columns=[
    # Edad
    'DAYS_BIRTH',

    # Montos base
    'AMT_CREDIT',
    'AMT_INCOME_TOTAL',
    'AMT_ANNUITY',

    # Scores individuales
    'EXT_SOURCE_1',
    'EXT_SOURCE_2',
    'EXT_SOURCE_3',

    # Familia
    'CNT_FAM_MEMBERS',
```

```

# Activos
'FLAG_OWN_CAR',
'FLAG_OWN_REALTY',

# Consultas buró total
'TOTAL_CONSULTAS_BURO',

# Componentes de ratios (solo MESES_CON_MORA está en df)
'MESES_CON_MORA',

# Variables binarias derivadas
'TIENE_IMPAGOS',
'ES_PRIMER_CREDITO',

# Créditos buró
'CANTIDAD_CREDITOS_BURO'
])

```

In [18]:

```

# =====
# RANDOM FOREST 2 – MODELO CON VARIABLES REDUCIDAS
# =====

print("=*80")
print("RANDOM FOREST 2 – MODELO CON VARIABLES REDUCIDAS")
print("=*80")

print(f"\nDataset después de drop: {df.shape}")
print(f"Variables eliminadas: 14")
print(f"Variables restantes: {len(df.columns)}")

=====

=====
RANDOM FOREST 2 – MODELO CON VARIABLES REDUCIDAS
=====

=====

Dataset después de drop: (307511, 26)
Variables eliminadas: 14
Variables restantes: 26

```

In [19]:

```

# =====
# PREPARACIÓN DE DATOS – MODELO 2
# =====

print("\n[2/15] Preparando datos para modelo 2...")

X2 = df.drop(['SK_ID_CURR', 'TARGET'], axis=1)
y2 = df['TARGET']

print(f"Variables predictoras: {X2.shape[1]}")
print(f"Variable objetivo: {y2.name}")

numeric_vars2 = X2.select_dtypes(include=[np.number]).columns.tolist()
categorical_vars2 = X2.select_dtypes(include=['object']).columns.tolist()

print(f"Variables numéricas: {len(numeric_vars2)}")
print(f"Variables categóricas: {len(categorical_vars2)}")

```

```

print(f"\nVariables en el modelo:")
for i, col in enumerate(X2.columns, 1):
    print(f" {i}. {col}")

```

[2/15] Preparando datos para modelo 2...
Variables predictoras: 24
Variable objetivo: TARGET
Variables numericas: 21
Variables categoricas: 3

Variables en el modelo:

1. EDAD_ANOS
2. SCORE_PROMEDIO
3. CREDIT_INCOME_RATIO
4. NAME_FAMILY_STATUS
5. CNT_CHILDREN
6. CODE_GENDER
7. NAME_EDUCATION_TYPE
8. INGRESO_PER_CAPITA
9. NUM_ACTIVOS
10. TOTAL_CREDITO_DISPONIBLE
11. TOTAL_CREDITO_OTORGADO
12. TOTAL_DEUDA_ACTUAL
13. MAX_DIAS_MORA
14. CREDITOS_ACTIVOS
15. CREDITOS_CERRADOS
16. PCT_MESES_MORA
17. CREDITOS_CON_IMPAGO
18. NUM_PRESTAMOS_PREVIOS
19. TASA_INTERES_PROMEDIO
20. PLAZO_PROMEDIO
21. MONTO_PROMEDIO_PREVIO
22. TOTAL_CREDITO_HISTORICO
23. RATIO_PAGO CUOTA
24. RATIO_PAGO_MINIMO_TC

In [20]:

```

# =====
# MANEJO DE VALORES FALTANTES - MODELO 2
# =====
print("\n[3/15] Analizando valores faltantes...")

missing_pct2 = (X2.isnull().sum() / len(X2) * 100).sort_values(ascending=False)
missing_vars2 = missing_pct2[missing_pct2 > 0]

if len(missing_vars2) > 0:
    print(f"Variables con valores faltantes: {len(missing_vars2)}")
    print(missing_vars2.head(10))

    print("\nImputando valores faltantes...")
    for col in numeric_vars2:
        if X2[col].isnull().sum() > 0:
            X2[col].fillna(X2[col].median(), inplace=True)

    for col in categorical_vars2:
        if X2[col].isnull().sum() > 0:
            X2[col].fillna(X2[col].mode()[0], inplace=True)

```

```
    print("Valores faltantes imputados")
else:
    print("No hay valores faltantes")
```

```
[3/15] Analizando valores faltantes...
Variables con valores faltantes: 5
TASA_INTERES_PROMEDIO      98.501192
RATIO_PAGO_MINIMO_TC       80.726218
PLAZO_PROMEDIO              5.485657
RATIO_PAGO_CUOTA            5.163718
SCORE_PROMEDIO               0.055933
dtype: float64
```

```
Imputando valores faltantes...
Valores faltantes imputados
```

In [21]:

```
# =====
# CODIFICACIÓN DE VARIABLES CATEGÓRICAS – MODELO 2
# =====
print("\n[4/15] Codificando variables categoricas...")

label_encoders2 = {}

if len(categorical_vars2) > 0:
    for col in categorical_vars2:
        le = LabelEncoder()
        X2[col] = le.fit_transform(X2[col].astype(str))
        label_encoders2[col] = le

    print(f"{len(categorical_vars2)} variables codificadas")
    print(f"Categorías: {categorical_vars2}")
else:
    print("No hay variables categoricas para codificar")
```

```
[4/15] Codificando variables categoricas...
```

3 variables codificadas

Categorías: ['NAME_FAMILY_STATUS', 'CODE_GENDER', 'NAME_EDUCATION_TYPE']

In [22]:

```
# =====
# DIVISIÓN TRAIN/TEST – MODELO 2
# =====
print("\n[5/15] Dividiendo datos en Train/Test...")

X2_train, X2_test, y2_train, y2_test = train_test_split(
    X2,
    y2,
    test_size=TEST_SIZE,
    random_state=RANDOM_STATE,
    stratify=y2
)

print(f"Train set: {X2_train.shape[0]} usuarios ({X2_train.shape[0]}/{len(X2)*2}%)")
print(f"Test set: {X2_test.shape[0]} usuarios ({X2_test.shape[0]}/{len(X2)*2}%)")
print(f"\nDistribucion TARGET en Train:")
print(f"  Clase 0 (Pago): {(y2_train==0).sum()}/{len(y2_train)}")
print(f"  Clase 1 (Default): {(y2_train==1).sum()}/{len(y2_train)}")
```

```
[5/15] Dividiendo datos en Train/Test...
```

```
Train set: 246,008 usuarios (80.0%)
```

```
Test set: 61,503 usuarios (20.0%)
```

```
Distribucion TARGET en Train:
```

```
Clase 0 (Pago): 226,148 (91.93%)
```

```
Clase 1 (Default): 19,860 (8.07%)
```

```
In [23]:
```

```
# =====
# ENTRENAMIENTO DEL MODELO 2
# =====
print("\n[7/15] Entrenando Random Forest 2 (variables reducidas)...")

model2 = RandomForestClassifier(
    n_estimators=N_ESTIMATORS,
    max_depth=MAX_DEPTH,
    min_samples_split=MIN_SAMPLES_SPLIT,
    min_samples_leaf=MIN_SAMPLES_LEAF,
    class_weight='balanced',
    random_state=RANDOM_STATE,
    n_jobs=-1,
    verbose=1
)

model2.fit(X2_train, y2_train)

print("\nModelo 2 entrenado exitosamente")
print(f"Numero de arboles: {model2.n_estimators}")
print(f"Profundidad maxima: {model2.max_depth}")
print(f"Min samples split: {model2.min_samples_split}")
print(f"Min samples leaf: {model2.min_samples_leaf}")
print(f"Clases: {model2.classes_}")
```

```
[7/15] Entrenando Random Forest 2 (variables reducidas)...
```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 4.0s
```

```
Modelo 2 entrenado exitosamente
```

```
Numero de arboles: 100
```

```
Profundidad maxima: 15
```

```
Min samples split: 10
```

```
Min samples leaf: 5
```

```
Clases: [0 1]
```

```
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 10.0s finished
```

```
In [24]:
```

```
# =====
# VALIDACIÓN CRUZADA – MODELO 2
# =====
print(f"\n[8/15] Validacion Cruzada ({CV_FOLDS}-fold en Train)...")

cv_scores2 = cross_val_score(
    model2, X2_train, y2_train,
    cv=CV_FOLDS,
    scoring='roc_auc',
    n_jobs=-1
)
```

```

print(f"ROC-AUC por fold: {cv_scores2}")
print(f"Media: {cv_scores2.mean():.4f} (+/- {cv_scores2.std():.4f})")

[8/15] Validacion Cruzada (5-fold en Train)...
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:  11.5s
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:  11.6s
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:  11.8s
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:  11.9s
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:  11.9s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  30.0s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  30.1s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  30.3s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  30.4s finished
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:  30.4s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.2s finished
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.2s finished
ROC-AUC por fold: [0.73840893 0.725804  0.73686891 0.73917831 0.74439579]
Media: 0.7369 (+/- 0.0061)

```

In [25]:

```

# =====
# PREDICCIONES - MODELO 2
# =====
print("\n[9/15] Generando predicciones...")

y2_train_pred = model2.predict(X2_train)
y2_train_proba = model2.predict_proba(X2_train)[:, 1]

```

```

y2_test_pred = model2.predict(X2_test)
y2_test_proba = model2.predict_proba(X2_test)[:, 1]

print("Predicciones generadas")

```

[9/15] Generando predicciones...

```

[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.4s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.1s
Predicciones generadas
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.4s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed:  0.1s finished

```

In [26]:

```

# =====
# EVALUACIÓN EN TRAIN – MODELO 2
# =====

print("\n[10/15] Evaluando modelo 2 en TRAIN SET...")
print("=*80")

train2_auc = roc_auc_score(y2_train, y2_train_proba)
train2_ap = average_precision_score(y2_train, y2_train_proba)

print(f"\nMetricas Globales:")
print(f"ROC-AUC Score: {train2_auc:.4f}")
print(f"Average Precision: {train2_ap:.4f}")

print(f"\nClassification Report:")
print(classification_report(y2_train, y2_train_pred,
                           target_names=['Pago (0)', 'Default (1)']))

print(f"\nConfusion Matrix:")
cm2_train = confusion_matrix(y2_train, y2_train_pred)
print(cm2_train)
print(f"\nInterpretacion:")
print(f" Verdaderos Negativos (TN): {cm2_train[0,0]} - Predijo 'Pago' y s")
print(f" Falsos Positivos (FP): {cm2_train[0,1]} - Predijo 'Default' pero")
print(f" Falsos Negativos (FN): {cm2_train[1,0]} - Predijo 'Pago' pero hi")
print(f" Verdaderos Positivos (TP): {cm2_train[1,1]} - Predijo 'Default'"

```

```
[10/15] Evaluando modelo 2 en TRAIN SET...
```

```
=====
```

Metricas Globales:

ROC-AUC Score: 0.9038

Average Precision: 0.5131

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Pago (0) | 0.98 | 0.86 | 0.91 | 226148 |
| Default (1) | 0.32 | 0.76 | 0.45 | 19860 |
| accuracy | | | 0.85 | 246008 |
| macro avg | 0.65 | 0.81 | 0.68 | 246008 |
| weighted avg | 0.92 | 0.85 | 0.88 | 246008 |

Confusion Matrix:

```
[[193755  32393]
 [ 4744 15116]]
```

Interpretacion:

Verdaderos Negativos (TN): 193,755 – Predijo 'Pago' y si pago

Falsos Positivos (FP): 32,393 – Predijo 'Default' pero pago

Falsos Negativos (FN): 4,744 – Predijo 'Pago' pero hizo default [CRITICO]

Verdaderos Positivos (TP): 15,116 – Predijo 'Default' y si hizo default

In [27]:

```
# =====
# EVALUACIÓN EN TEST – MODELO 2
# =====

print("\n[11/15] Evaluando modelo 2 en TEST SET (metricas definitivas)...")
print("=*80)

test2_auc = roc_auc_score(y2_test, y2_test_proba)
test2_ap = average_precision_score(y2_test, y2_test_proba)

print(f"\nMetricas Globales:")
print(f"ROC-AUC Score: {test2_auc:.4f}")
print(f"Average Precision: {test2_ap:.4f}")

print(f"\nClassification Report:")
print(classification_report(y2_test, y2_test_pred,
                           target_names=['Pago (0)', 'Default (1)']))

print(f"\nConfusion Matrix:")
cm2_test = confusion_matrix(y2_test, y2_test_pred)
print(cm2_test)
print(f"\nInterpretacion:")
print(f" Verdaderos Negativos (TN): {cm2_test[0,0]} – Predijo 'Pago' y si
print(f" Falsos Positivos (FP): {cm2_test[0,1]} – Predijo 'Default' pero
print(f" Falsos Negativos (FN): {cm2_test[1,0]} – Predijo 'Pago' pero fiz
print(f" Verdaderos Positivos (TP): {cm2_test[1,1]} – Predijo 'Default' y
```

[11/15] Evaluando modelo 2 en TEST SET (metricas definitivas)...

```
=====
```

Metricas Globales:

ROC-AUC Score: 0.7407

Average Precision: 0.2133

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Pago (0) | 0.95 | 0.84 | 0.89 | 56538 |
| Default (1) | 0.20 | 0.46 | 0.28 | 4965 |
| accuracy | | | 0.81 | 61503 |
| macro avg | 0.58 | 0.65 | 0.59 | 61503 |
| weighted avg | 0.89 | 0.81 | 0.84 | 61503 |

Confusion Matrix:

```
[[47558 8980]
 [ 2661 2304]]
```

Interpretacion:

Verdaderos Negativos (TN): 47,558 – Predijo 'Pago' y si pago

Falsos Positivos (FP): 8,980 – Predijo 'Default' pero pago

Falsos Negativos (FN): 2,661 – Predijo 'Pago' pero hizo default [CRITICO]

Verdaderos Positivos (TP): 2,304 – Predijo 'Default' y si hizo default

In [28]:

```
# =====
# COMPARACIÓN TRAIN VS TEST - MODELO 2
# =====

print("\n[12/15] Comparando TRAIN vs TEST (analisis de overfitting)...")
print("=*80)

print(f"\n{'Metrica':<25} {'Train':<12} {'Test':<12} {'Diferencia':<12}")
print("-" * 65)
print(f"{'ROC-AUC':<25} {train2_auc:<12.4f} {test2_auc:<12.4f} {abs(train2_a
print(f"{'Average Precision':<25} {train2_ap:<12.4f} {test2_ap:<12.4f} {abs(t

diff_auc2 = abs(train2_auc - test2_auc)
if diff_auc2 < 0.02:
    print("\nExcelente generalizacion (diferencia < 2%)")
elif diff_auc2 < 0.05:
    print("\nBuena generalizacion (diferencia < 5%)")
elif diff_auc2 < 0.10:
    print("\nPosible ligero overfitting (diferencia 5-10%)")
else:
    print("\nOverfitting detectado (diferencia > 10%)")
```

[12/15] Comparando TRAIN vs TEST (analisis de overfitting)...

=====

| Metrica | Train | Test | Diferencia |
|-------------------|--------|--------|------------|
| ROC-AUC | 0.9038 | 0.7407 | 0.1631 |
| Average Precision | 0.5131 | 0.2133 | 0.2998 |

Overfitting detectado (diferencia > 10%)

In [29]:

```
# =====
# IMPORTANCIA DE VARIABLES – MODELO 2
# =====
print("\n[13/15] Calculando importancia de variables...")
print("=*80)

feature_importance2 = pd.DataFrame({
    'Variable': X2.columns,
    'Importancia': model2.feature_importances_
}).sort_values('Importancia', ascending=False)

print("\nTop 20 variables mas importantes para predecir default:\n")
print(feature_importance2.head(20).to_string(index=False))

print("\nInterpretacion de importancia (Gini Importance):")
print(" - Problema de CLASIFICACIÓN BINARIA: predecir 0 (Pago) o 1 (Default")
print(" - Gini mide cuanto reduce cada variable la impureza en los splits")
print(" - Valores mas altos = mayor capacidad para separar las clases")
print(" - Basado en reduccion promedio de impureza en los 100 arboles")
```

[13/15] Calculando importancia de variables...

```
=====
```

Top 20 variables mas importantes para predecir default:

| Variable | Importancia |
|--------------------------|-------------|
| SCORE_PROMEDIO | 0.321122 |
| EDAD_ANOS | 0.082152 |
| CREDIT_INCOME_RATIO | 0.054617 |
| RATIO_PAGO_CUOTA | 0.053444 |
| TOTAL_CREDITO_OTORGADO | 0.050914 |
| TOTAL_CREDITO_HISTORICO | 0.047529 |
| MONTO_PROMEDIO_PREVIO | 0.046952 |
| PLAZO_PROMEDIO | 0.045953 |
| TOTAL_DEUDA_ACTUAL | 0.040421 |
| INGRESO_PER_CAPITA | 0.038446 |
| CREDITOS_CERRADOS | 0.030508 |
| RATIO_PAGO_MINIMO_TC | 0.026664 |
| NAME_EDUCATION_TYPE | 0.025050 |
| NUM_PRESTAMOS_PREVIOS | 0.024877 |
| CREDITOS_ACTIVOS | 0.023181 |
| TOTAL_CREDITO_DISPONIBLE | 0.019034 |
| CODE_GENDER | 0.017551 |
| NAME_FAMILY_STATUS | 0.012794 |
| PCT_MESES_MORA | 0.012441 |
| NUM_ACTIVOS | 0.010104 |

Interpretacion de importancia (Gini Importance):

- Problema de CLASIFICACIÓN BINARIA: predecir 0 (Pago) o 1 (Default)
- Gini mide cuanto reduce cada variable la impureza en los splits
- Valores mas altos = mayor capacidad para separar las clases
- Basado en reduccion promedio de impureza en los 100 arboles

In [30]:

```
# =====
# VISUALIZACIONES - MODELO 2
# =====
print("\n[14/15] Generando visualizaciones...")

# Calcular deciles
test_results2 = pd.DataFrame({
    'y_true': y2_test,
    'y_proba': y2_test_proba
})
test_results2['decil'] = pd.qcut(test_results2['y_proba'], q=10, labels=False)

# Calcular TPR por decil
decil_stats2 = test_results2.groupby('decil').agg({
    'y_true': ['sum', 'count', 'mean']
}).reset_index()
decil_stats2.columns = ['decil', 'positivos', 'total', 'tpr']
decil_stats2 = decil_stats2.sort_values('decil', ascending=False)

# Crear gráficas
fig, axes = plt.subplots(2, 2, figsize=(16, 14))
```

```

# ROC Curve
fpr2, tpr2, _ = roc_curve(y2_test, y2_test_proba)
axes[0, 0].plot(fpr2, tpr2, linewidth=2, color='darkorange', label=f'ROC (All')
axes[0, 0].plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random')
axes[0, 0].set_xlabel('False Positive Rate', fontsize=12)
axes[0, 0].set_ylabel('True Positive Rate', fontsize=12)
axes[0, 0].set_title('ROC Curve – Test Set (RF2 – Variables Reducidas)', fontw
axes[0, 0].legend()
axes[0, 0].grid(alpha=0.3)

# Precision-Recall Curve
precision2, recall2, _ = precision_recall_curve(y2_test, y2_test_proba)
sort_idx2 = np.argsort(recall2)
recall_sorted2 = recall2[sort_idx2]
precision_sorted2 = precision2[sort_idx2]

axes[0, 1].plot(recall_sorted2, precision_sorted2, linewidth=2, color='darkorange')
axes[0, 1].axhline(y=y2_test.mean(), color='gray', linestyle='--', linewidth=1)
axes[0, 1].set_xlabel('Recall (TPR)', fontsize=12)
axes[0, 1].set_ylabel('Precision', fontsize=12)
axes[0, 1].set_title('Precision-Recall Curve – Test Set', fontsize=14, fontweight='bold')
axes[0, 1].set_xlim([0, 1])
axes[0, 1].set_ylim([0, 1])
axes[0, 1].legend()
axes[0, 1].grid(alpha=0.3)

# TPR por Decil
axes[1, 0].bar(decil_stats2['decil'], decil_stats2['tpr'], color='darkorange')
axes[1, 0].set_xlabel('Decil de Score (10 = Mayor Riesgo)', fontsize=12)
axes[1, 0].set_ylabel('True Positive Rate (TPR)', fontsize=12)
axes[1, 0].set_title('TPR por Decil de Probabilidad', fontsize=14, fontweight='bold')
axes[1, 0].set_xticks(range(1, 11))
axes[1, 0].grid(alpha=0.3, axis='y')
axes[1, 0].set_ylim(0, 1)

for i, row in decil_stats2.iterrows():
    axes[1, 0].text(row['decil'], row['tpr'] + 0.02, f'{row["tpr"]:.2f}', ha='center', va='bottom', fontsize=10)

# Feature Importance (Top 15)
top_features2 = feature_importance2.head(15)
axes[1, 1].barh(range(len(top_features2)), top_features2['Importancia'].values)
axes[1, 1].set_yticks(range(len(top_features2)))
axes[1, 1].set_yticklabels(top_features2['Variable'].values)
axes[1, 1].set_xlabel('Importancia (Gini)', fontsize=12)
axes[1, 1].set_title('Top 15 Variables Mas Importantes (RF2)', fontsize=14, fontweight='bold')
axes[1, 1].invert_yaxis()
axes[1, 1].grid(alpha=0.3, axis='x')

plt.tight_layout()
plt.savefig(f'{OUTPUT_DIR}/random_forest2_curves.png', dpi=300, bbox_inches='tight')
print(f"Graficas guardadas: {OUTPUT_DIR}/random_forest2_curves.png")

plt.show()

# Imprimir tabla de deciles

```

```

print("\nTabla de TPR por Decil:")
print(decil_stats2.to_string(index=False))

```

[14/15] Generando visualizaciones...

Graficas guardadas: ../../data/random-forest-output/random_forest2_curves.png

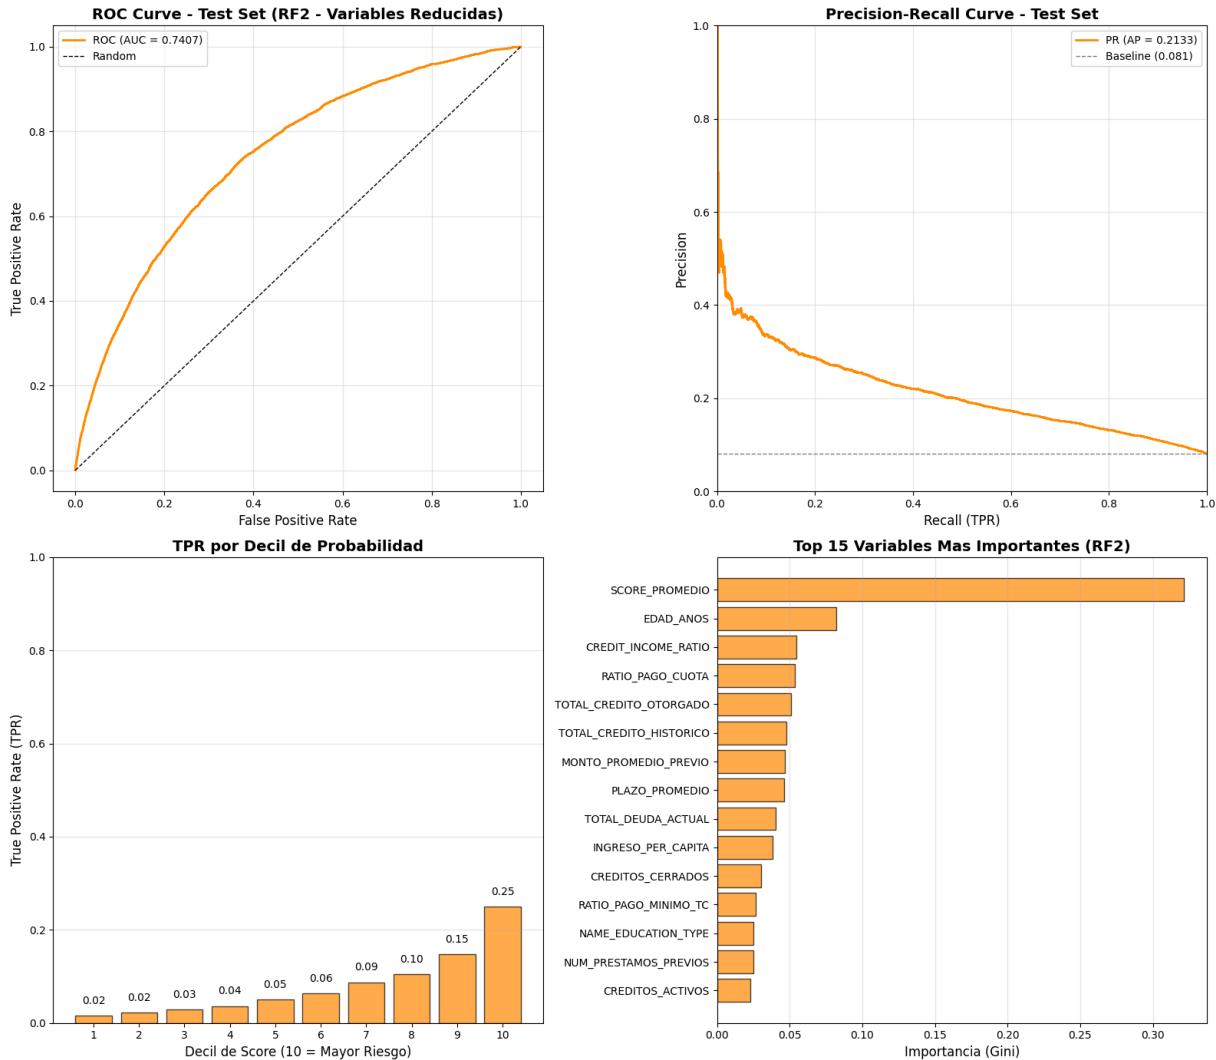


Tabla de TPR por Decil:

| decil | positivos | total | tpr |
|-------|-----------|-------|----------|
| 10 | 1533 | 6151 | 0.249228 |
| 9 | 912 | 6150 | 0.148293 |
| 8 | 644 | 6150 | 0.104715 |
| 7 | 538 | 6150 | 0.087480 |
| 6 | 391 | 6150 | 0.063577 |
| 5 | 313 | 6151 | 0.050886 |
| 4 | 224 | 6150 | 0.036423 |
| 3 | 178 | 6150 | 0.028943 |
| 2 | 138 | 6150 | 0.022439 |
| 1 | 94 | 6151 | 0.015282 |

In [31]:

```

# =====
# GUARDAR MODELO Y RESULTADOS – MODELO 2
# =====
print("\n[15/15] Guardando modelo 2 y resultados...")

with open(f'{OUTPUT_DIR}/random_forest2_model.pkl', 'wb') as f:

```

```

    pickle.dump(model2, f)

    with open(f'{OUTPUT_DIR}/label_encoders_rf2.pkl', 'wb') as f:
        pickle.dump(label_encoders2, f)

    feature_importance2.to_csv(f'{OUTPUT_DIR}/feature_importance_rf2.csv', index=False)

    test_predictions2 = pd.DataFrame({
        'SK_ID_CURR': df.iloc[X2_test.index]['SK_ID_CURR'].values,
        'TARGET_Real': y2_test.values,
        'TARGET_Predicho': y2_test_pred,
        'Probabilidad_Default': y2_test_proba
    })
    test_predictions2.to_csv(f'{OUTPUT_DIR}/test_predictions_rf2.csv', index=False)

    print(f"Archivos guardados en: {OUTPUT_DIR}/")
    print(" - random_forest2_model.pkl")
    print(" - label_encoders_rf2.pkl")
    print(" - feature_importance_rf2.csv")
    print(" - test_predictions_rf2.csv")
    print(" - random_forest2_curves.png")

```

[15/15] Guardando modelo 2 y resultados...

Archivos guardados en: ../../data/random-forest-output/

- random_forest2_model.pkl
- label_encoders_rf2.pkl
- feature_importance_rf2.csv
- test_predictions_rf2.csv
- random_forest2_curves.png

In []:

```

# =====
# COMPARACIÓN MODELO 1 VS MODELO 2
# =====

print("\n" + "="*80)
print("COMPARACIÓN: MODELO 1 (39 vars) VS MODELO 2 (25 vars)")
print("="*80)

print(f"\n{'Metrica':<25} {'RF1 (39 vars)':<15} {'RF2 (25 vars)':<15} {'Diferencia AUC':<15}")
print("-" * 70)
print(f"{'ROC-AUC (Test)':<25} {test_auc:<15.4f} {test2_auc:<15.4f} {test2_auc - test_auc:<15.4f}")
print(f"{'Average Precision (Test)':<25} {test_ap:<15.4f} {test2_ap:<15.4f} {test2_ap - test_ap:<15.4f}")
print(f"{'ROC-AUC (Train)':<25} {train_auc:<15.4f} {train2_auc:<15.4f} {train2_auc - train_auc:<15.4f}")
print(f"{'Overfitting (Train-Test)':<25} {train_auc - test_auc:<15.4f} {train2_auc - test2_auc:<15.4f}")

print(f"\nNúmero de variables:")
print(f"  Modelo 1: {X.shape[1]} variables")
print(f"  Modelo 2: {X2.shape[1]} variables")
print(f"  Reducción: {X.shape[1] - X2.shape[1]} variables eliminadas ({(X.shape[1] - X2.shape[1]) / X.shape[1] * 100:.2f}%)")

# Conclusión
if test2_auc >= test_auc:
    print(f"\n✓ Modelo 2 tiene MEJOR o IGUAL rendimiento con menos variables")
    print(f"  Esto sugiere que las variables eliminadas eran redundantes")
else:
    print(f"\nx Modelo 2 tiene PEOR rendimiento")
    print(f"  Diferencia de AUC: {test_auc - test2_auc:.4f}")

```

```

print("\n" + "="*80)
print("ANALISIS COMPLETADO – RANDOM FOREST 2 (VARIABLES REDUCIDAS)")
print("="*80)
=====
```

```
=====
=====
COMPARACIÓN: MODELO 1 (39 vars) VS MODELO 2 (25 vars)
=====
```

| Metrica | RF1 (39 vars) | RF2 (25 vars) | Diferencia |
|--------------------------|---------------|---------------|------------|
| ROC-AUC (Test) | 0.7451 | 0.7407 | -0.0044 |
| Average Precision (Test) | 0.2207 | 0.2133 | -0.0074 |
| ROC-AUC (Train) | 0.9222 | 0.9038 | -0.0184 |
| Overfitting (Train-Test) | 0.1771 | 0.1631 | |

Numero de variables:

Modelo 1: 39 variables

Modelo 2: 24 variables

Reducción: 15 variables eliminadas (38.5%)

✗ Modelo 2 tiene PEOR rendimiento

Diferencia de AUC: 0.0044

```
=====
=====
ANALISIS COMPLETADO – RANDOM FOREST 2 (VARIABLES REDUCIDAS)
=====
```

```
=====
```

Optimización de hiperparámetros

In [33]:

```

# =====
# RANDOM FOREST 3 – OPTIMIZACIÓN DE HIPERPARÁMETROS (MODELO 2)
# =====
# Usaremos RandomizedSearchCV para encontrar los mejores hiperparámetros
# para el modelo con variables reducidas (24 variables)

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint, uniform

print("*80)
print("RANDOM FOREST 3 – BÚSQUEDA DE HIPERPARÁMETROS ÓPTIMOS")
print("*80)
print(f"\nUsando el dataset de {X2.shape[1]} variables (Modelo 2)")
print(f"Datos de entrenamiento: {X2_train.shape[0]}, usuarios")
print(f"Datos de prueba: {X2_test.shape[0]}, usuarios")
```

```
=====
===== RANDOM FOREST 3 – BÚSQUEDA DE HIPERPARÁMETROS ÓPTIMOS =====
```

Usando el dataset de 24 variables (Modelo 2)
Datos de entrenamiento: 246,008 usuarios
Datos de prueba: 61,503 usuarios

```
In [34]: # -----
# DEFINICIÓN DEL ESPACIO DE BÚSQUEDA
# -----
print("\n[1/5] Definiendo espacio de búsqueda de hiperparámetros...")

# Espacio de hiperparámetros a explorar
param_distributions = {
    # Número de árboles – más árboles = mejor generalización pero más tiempo
    'n_estimators': [50, 100, 150, 200, 300],

    # Profundidad máxima – controla overfitting
    # Valores más bajos = menos overfitting
    'max_depth': [5, 8, 10, 12, 15, 20, None],

    # Mínimo de muestras para split – valores más altos = menos overfitting
    'min_samples_split': [2, 5, 10, 15, 20, 30, 50],

    # Mínimo de muestras en hoja – valores más altos = menos overfitting
    'min_samples_leaf': [1, 2, 5, 10, 15, 20, 30],

    # Número de features a considerar en cada split
    'max_features': ['sqrt', 'log2', 0.3, 0.5, 0.7],

    # Bootstrap samples
    'bootstrap': [True, False],
}

# Criterio de split
'criterion': ['gini', 'entropy']
}

# Calcular número total de combinaciones
total_combinations = 1
for key, values in param_distributions.items():
    total_combinations *= len(values)

print(f"\nHiperparámetros a optimizar:")
for param, values in param_distributions.items():
    print(f" - {param}: {values}")

print(f"\nCombinaciones totales posibles: {total_combinations:,}")
print(f"Iteraciones a probar: 50 (RandomizedSearchCV)")
```

[1/5] Definiendo espacio de búsqueda de hiperparámetros...

Hiperparámetros a optimizar:

- n_estimators: [50, 100, 150, 200, 300]
- max_depth: [5, 8, 10, 12, 15, 20, None]
- min_samples_split: [2, 5, 10, 15, 20, 30, 50]
- min_samples_leaf: [1, 2, 5, 10, 15, 20, 30]
- max_features: ['sqrt', 'log2', 0.3, 0.5, 0.7]
- bootstrap: [True, False]
- criterion: ['gini', 'entropy']

Combinaciones totales posibles: 34,300

Iteraciones a probar: 50 (RandomizedSearchCV)

```
In [36]: # =====#
# BÚSQUEDA DE HIPERPARÁMETROS CON RANDOMIZED SEARCH CV
# =====#
print("\n[2/5] Ejecutando RandomizedSearchCV...")
print("NOTA: Este proceso puede tardar varios minutos...")

# Modelo base
rf_base = RandomForestClassifier(
    class_weight='balanced',
    random_state=RANDOM_STATE,
    n_jobs=1
)

# RandomizedSearchCV con validación cruzada
# IMPORTANTE: n_jobs=1 para evitar errores de multiprocessing en Jupyter notebook
random_search = RandomizedSearchCV(
    estimator=rf_base,
    param_distributions=param_distributions,
    n_iter=30, # Reducido a 30 para acelerar
    cv=3, # 3-fold CV para acelerar
    scoring='roc_auc',
    random_state=RANDOM_STATE,
    n_jobs=1, # Evita FileNotFoundError en notebooks
    verbose=2,
    return_train_score=True
)

# Ejecutar búsqueda
import time
start_time = time.time()

random_search.fit(X2_train, y2_train)

elapsed_time = time.time() - start_time
print(f"\nBúsqueda completada en {elapsed_time/60:.2f} minutos!")
```

```
[2/5] Ejecutando RandomizedSearchCV...
NOTA: Este proceso puede tardar varios minutos...
Fitting 3 folds for each of 30 candidates, totalling 90 fits
[CV] END bootstrap=True, criterion=entropy, max_depth=20, max_features=0.7,
min_samples_leaf=10, min_samples_split=10, n_estimators=50; total time= 13.
1s
[CV] END bootstrap=True, criterion=entropy, max_depth=20, max_features=0.7,
min_samples_leaf=10, min_samples_split=10, n_estimators=50; total time= 12.
2s
[CV] END bootstrap=True, criterion=entropy, max_depth=20, max_features=0.7,
min_samples_leaf=10, min_samples_split=10, n_estimators=50; total time= 12.
0s
[CV] END bootstrap=True, criterion=gini, max_depth=5, max_features=0.5, min_
samples_leaf=10, min_samples_split=20, n_estimators=50; total time= 2.5s
[CV] END bootstrap=True, criterion=gini, max_depth=5, max_features=0.5, min_
samples_leaf=10, min_samples_split=20, n_estimators=50; total time= 2.8s
[CV] END bootstrap=True, criterion=gini, max_depth=5, max_features=0.5, min_
samples_leaf=10, min_samples_split=20, n_estimators=50; total time= 2.7s
[CV] END bootstrap=True, criterion=entropy, max_depth=10, max_features=log2,
min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 14.
1s
[CV] END bootstrap=True, criterion=entropy, max_depth=10, max_features=log2,
min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 12.
7s
[CV] END bootstrap=True, criterion=entropy, max_depth=10, max_features=log2,
min_samples_leaf=1, min_samples_split=10, n_estimators=300; total time= 12.
8s
[CV] END bootstrap=True, criterion=gini, max_depth=20, max_features=sqrt, mi
n_samples_leaf=15, min_samples_split=2, n_estimators=50; total time= 2.7s
[CV] END bootstrap=True, criterion=gini, max_depth=20, max_features=sqrt, mi
n_samples_leaf=15, min_samples_split=2, n_estimators=50; total time= 2.7s
[CV] END bootstrap=True, criterion=gini, max_depth=20, max_features=sqrt, mi
n_samples_leaf=15, min_samples_split=2, n_estimators=50; total time= 2.9s
[CV] END bootstrap=True, criterion=entropy, max_depth=None, max_features=0.
5, min_samples_leaf=20, min_samples_split=15, n_estimators=50; total time=
8.5s
[CV] END bootstrap=True, criterion=entropy, max_depth=None, max_features=0.
5, min_samples_leaf=20, min_samples_split=15, n_estimators=50; total time=
8.4s
[CV] END bootstrap=True, criterion=entropy, max_depth=None, max_features=0.
5, min_samples_leaf=20, min_samples_split=15, n_estimators=50; total time=
8.8s
[CV] END bootstrap=False, criterion=gini, max_depth=12, max_features=0.7, mi
n_samples_leaf=15, min_samples_split=15, n_estimators=150; total time= 35.5
s
[CV] END bootstrap=False, criterion=gini, max_depth=12, max_features=0.7, mi
n_samples_leaf=15, min_samples_split=15, n_estimators=150; total time= 34.0
s
[CV] END bootstrap=False, criterion=gini, max_depth=12, max_features=0.7, mi
n_samples_leaf=15, min_samples_split=15, n_estimators=150; total time= 34.1
s
[CV] END bootstrap=True, criterion=entropy, max_depth=None, max_features=sqr
t, min_samples_leaf=5, min_samples_split=30, n_estimators=200; total time=
14.6s
[CV] END bootstrap=True, criterion=entropy, max_depth=None, max_features=sqr
t, min_samples_leaf=5, min_samples_split=30, n_estimators=200; total time=
```

14.7s

[CV] END bootstrap=True, criterion=entropy, max_depth=None, max_features=sqrt, min_samples_leaf=5, min_samples_split=30, n_estimators=200; total time= 15.8s

[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=log2, min_samples_leaf=30, min_samples_split=5, n_estimators=50; total time= 2.1s

[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=log2, min_samples_leaf=30, min_samples_split=5, n_estimators=50; total time= 1.8s

[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=log2, min_samples_leaf=30, min_samples_split=5, n_estimators=50; total time= 1.8s

[CV] END bootstrap=True, criterion=gini, max_depth=5, max_features=0.5, min_samples_leaf=1, min_samples_split=50, n_estimators=300; total time= 17.3s

[CV] END bootstrap=True, criterion=gini, max_depth=5, max_features=0.5, min_samples_leaf=1, min_samples_split=50, n_estimators=300; total time= 16.4s

[CV] END bootstrap=True, criterion=gini, max_depth=5, max_features=0.5, min_samples_leaf=1, min_samples_split=50, n_estimators=300; total time= 16.8s

[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=0.7, min_samples_leaf=30, min_samples_split=15, n_estimators=200; total time= 22.3s

[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=0.7, min_samples_leaf=30, min_samples_split=15, n_estimators=200; total time= 22.2s

[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=0.7, min_samples_leaf=30, min_samples_split=15, n_estimators=200; total time= 21.8s

[CV] END bootstrap=True, criterion=gini, max_depth=15, max_features=log2, min_samples_leaf=15, min_samples_split=30, n_estimators=100; total time= 5.0s

[CV] END bootstrap=True, criterion=gini, max_depth=15, max_features=log2, min_samples_leaf=15, min_samples_split=30, n_estimators=100; total time= 5.1s

[CV] END bootstrap=True, criterion=gini, max_depth=15, max_features=log2, min_samples_leaf=15, min_samples_split=30, n_estimators=100; total time= 5.2s

[CV] END bootstrap=False, criterion=gini, max_depth=5, max_features=log2, min_samples_leaf=15, min_samples_split=50, n_estimators=200; total time= 5.8s

[CV] END bootstrap=False, criterion=gini, max_depth=5, max_features=log2, min_samples_leaf=15, min_samples_split=50, n_estimators=200; total time= 5.9s

[CV] END bootstrap=False, criterion=gini, max_depth=5, max_features=log2, min_samples_leaf=15, min_samples_split=50, n_estimators=200; total time= 5.9s

[CV] END bootstrap=False, criterion=gini, max_depth=10, max_features=sqrt, min_samples_leaf=15, min_samples_split=30, n_estimators=300; total time= 17.2s

[CV] END bootstrap=False, criterion=gini, max_depth=10, max_features=sqrt, min_samples_leaf=15, min_samples_split=30, n_estimators=300; total time= 17.2s

[CV] END bootstrap=False, criterion=gini, max_depth=10, max_features=sqrt, min_samples_leaf=15, min_samples_split=30, n_estimators=300; total time= 17.3s

[CV] END bootstrap=False, criterion=entropy, max_depth=10, max_features=0.3, min_samples_leaf=1, min_samples_split=30, n_estimators=200; total time= 22.2s

[CV] END bootstrap=False, criterion=entropy, max_depth=10, max_features=0.3, min_samples_leaf=1, min_samples_split=30, n_estimators=200; total time= 22.8s

[CV] END bootstrap=False, criterion=entropy, max_depth=10, max_features=0.3,


```
samples_leaf=20, min_samples_split=5, n_estimators=300; total time= 16.3s
[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=0.3, min_
samples_leaf=20, min_samples_split=5, n_estimators=300; total time= 16.6s
[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=sqrt, min_
samples_leaf=2, min_samples_split=5, n_estimators=150; total time= 5.2s
[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=sqrt, min_
samples_leaf=2, min_samples_split=5, n_estimators=150; total time= 5.1s
[CV] END bootstrap=True, criterion=gini, max_depth=8, max_features=sqrt, min_
samples_leaf=2, min_samples_split=5, n_estimators=150; total time= 5.1s
[CV] END bootstrap=False, criterion=entropy, max_depth=15, max_features=0.5,
min_samples_leaf=20, min_samples_split=15, n_estimators=100; total time= 2
6.4s
[CV] END bootstrap=False, criterion=entropy, max_depth=15, max_features=0.5,
min_samples_leaf=20, min_samples_split=15, n_estimators=100; total time= 2
7.2s
[CV] END bootstrap=False, criterion=entropy, max_depth=15, max_features=0.5,
min_samples_leaf=20, min_samples_split=15, n_estimators=100; total time= 2
8.4s
[CV] END bootstrap=True, criterion=entropy, max_depth=10, max_features=log2,
min_samples_leaf=10, min_samples_split=15, n_estimators=300; total time= 1
5.1s
[CV] END bootstrap=True, criterion=entropy, max_depth=10, max_features=log2,
min_samples_leaf=10, min_samples_split=15, n_estimators=300; total time= 1
4.8s
[CV] END bootstrap=True, criterion=entropy, max_depth=10, max_features=log2,
min_samples_leaf=10, min_samples_split=15, n_estimators=300; total time= 1
4.6s
[CV] END bootstrap=True, criterion=gini, max_depth=10, max_features=0.7, min_
samples_leaf=10, min_samples_split=20, n_estimators=100; total time= 15.2s
[CV] END bootstrap=True, criterion=gini, max_depth=10, max_features=0.7, min_
samples_leaf=10, min_samples_split=20, n_estimators=100; total time= 15.4s
[CV] END bootstrap=True, criterion=gini, max_depth=10, max_features=0.7, min_
samples_leaf=10, min_samples_split=20, n_estimators=100; total time= 16.1s
[CV] END bootstrap=True, criterion=gini, max_depth=12, max_features=sqrt, mi
n_samples_leaf=30, min_samples_split=5, n_estimators=50; total time= 2.6s
[CV] END bootstrap=True, criterion=gini, max_depth=12, max_features=sqrt, mi
n_samples_leaf=30, min_samples_split=5, n_estimators=50; total time= 2.5s
[CV] END bootstrap=True, criterion=gini, max_depth=12, max_features=sqrt, mi
n_samples_leaf=30, min_samples_split=5, n_estimators=50; total time= 2.7s
[CV] END bootstrap=False, criterion=entropy, max_depth=15, max_features=sqr
t, min_samples_leaf=10, min_samples_split=10, n_estimators=50; total time=
5.1s
[CV] END bootstrap=False, criterion=entropy, max_depth=15, max_features=sqr
t, min_samples_leaf=10, min_samples_split=10, n_estimators=50; total time=
5.1s
[CV] END bootstrap=False, criterion=entropy, max_depth=15, max_features=sqr
t, min_samples_leaf=10, min_samples_split=10, n_estimators=50; total time=
5.1s
[CV] END bootstrap=True, criterion=entropy, max_depth=15, max_features=0.7,
min_samples_leaf=2, min_samples_split=10, n_estimators=150; total time= 34.
6s
[CV] END bootstrap=True, criterion=entropy, max_depth=15, max_features=0.7,
min_samples_leaf=2, min_samples_split=10, n_estimators=150; total time= 37.
5s
[CV] END bootstrap=True, criterion=entropy, max_depth=15, max_features=0.7,
min_samples_leaf=2, min_samples_split=10, n_estimators=150; total time= 38.
```

5s

¡Búsqueda completada en 22.75 minutos!

```
In [48]: # =====
# MEJORES HIPERPARÁMETROS ENCONTRADOS
# =====
print("\n[3/5] Mejores hiperparámetros encontrados:")
print("=*80)

best_params = random_search.best_params_
best_score = random_search.best_score_

print(f"\nMejor ROC-AUC en CV: {best_score:.4f}")
print(f"\nHiperparámetros óptimos:")
for param, value in best_params.items():
    print(f"  {param}: {value}")

# Mostrar top 10 combinaciones
print("\n" + "*80)
print("Top 10 mejores combinaciones de hiperparámetros:")
print("*80)

cv_results = pd.DataFrame(random_search.cv_results_)
cv_results_sorted = cv_results.sort_values('rank_test_score').head(10)

for idx, row in cv_results_sorted.iterrows():
    print(f"\nRank {row['rank_test_score']:.0f}: ROC-AUC = {row['mean_test_s
    print(f"  Train AUC: {row['mean_train_score']:.4f}, Overfitting: {row['n
    params = row['params']
    for p, v in params.items():
        print(f"    {p}: {v}")
```

[3/5] Mejores hiperparámetros encontrados:

```
=====
```

====
Mejor ROC-AUC en CV: 0.7448

Hiperparámetros óptimos:

```
n_estimators: 200  
min_samples_split: 30  
min_samples_leaf: 5  
max_features: sqrt  
max_depth: None  
criterion: entropy  
bootstrap: True
```

```
=====
```

====

Top 10 mejores combinaciones de hiperparámetros:

```
=====
```

====

Rank 1: ROC-AUC = 0.7448 (+/- 0.0020)

```
Train AUC: 0.9909, Overfitting: 0.2461  
n_estimators: 200  
min_samples_split: 30  
min_samples_leaf: 5  
max_features: sqrt  
max_depth: None  
criterion: entropy  
bootstrap: True
```

Rank 2: ROC-AUC = 0.7432 (+/- 0.0014)

```
Train AUC: 0.8241, Overfitting: 0.0809  
n_estimators: 50  
min_samples_split: 5  
min_samples_leaf: 30  
max_features: sqrt  
max_depth: 12  
criterion: gini  
bootstrap: True
```

Rank 3: ROC-AUC = 0.7429 (+/- 0.0015)

```
Train AUC: 0.8944, Overfitting: 0.1515  
n_estimators: 100  
min_samples_split: 30  
min_samples_leaf: 15  
max_features: log2  
max_depth: 15  
criterion: gini  
bootstrap: True
```

Rank 4: ROC-AUC = 0.7429 (+/- 0.0018)

```
Train AUC: 0.8136, Overfitting: 0.0707  
n_estimators: 300  
min_samples_split: 30  
min_samples_leaf: 15
```

```
max_features: sqrt
max_depth: 10
criterion: gini
bootstrap: False

Rank 5: ROC-AUC = 0.7429 (+/- 0.0018)
Train AUC: 0.8015, Overfitting: 0.0586
n_estimators: 300
min_samples_split: 15
min_samples_leaf: 10
max_features: log2
max_depth: 10
criterion: entropy
bootstrap: True

Rank 6: ROC-AUC = 0.7421 (+/- 0.0018)
Train AUC: 0.8049, Overfitting: 0.0628
n_estimators: 300
min_samples_split: 10
min_samples_leaf: 1
max_features: log2
max_depth: 10
criterion: entropy
bootstrap: True

Rank 7: ROC-AUC = 0.7418 (+/- 0.0020)
Train AUC: 0.7731, Overfitting: 0.0313
n_estimators: 300
min_samples_split: 5
min_samples_leaf: 20
max_features: 0.3
max_depth: 8
criterion: gini
bootstrap: True

Rank 8: ROC-AUC = 0.7415 (+/- 0.0018)
Train AUC: 0.8078, Overfitting: 0.0664
n_estimators: 150
min_samples_split: 15
min_samples_leaf: 2
max_features: log2
max_depth: 10
criterion: gini
bootstrap: True

Rank 9: ROC-AUC = 0.7415 (+/- 0.0016)
Train AUC: 0.9424, Overfitting: 0.2009
n_estimators: 50
min_samples_split: 2
min_samples_leaf: 15
max_features: sqrt
max_depth: 20
criterion: gini
bootstrap: True

Rank 10: ROC-AUC = 0.7414 (+/- 0.0025)
```

```
Train AUC: 0.8172, Overfitting: 0.0758
n_estimators: 200
min_samples_split: 30
min_samples_leaf: 1
max_features: 0.3
max_depth: 10
criterion: entropy
bootstrap: False
```

```
In [41]: # =====#
# EVALUACIÓN DEL MODELO OPTIMIZADO
# =====#
print("\n[4/5] Evaluando modelo con hiperparámetros óptimos...")
print("*80")

# El mejor modelo ya está entrenado
best_model = random_search.best_estimator_

# Predicciones
y3_train_pred = best_model.predict(X2_train)
y3_train_proba = best_model.predict_proba(X2_train)[:, 1]

y3_test_pred = best_model.predict(X2_test)
y3_test_proba = best_model.predict_proba(X2_test)[:, 1]

# Métricas en Train
train3_auc = roc_auc_score(y2_train, y3_train_proba)
train3_ap = average_precision_score(y2_train, y3_train_proba)

# Métricas en Test
test3_auc = roc_auc_score(y2_test, y3_test_proba)
test3_ap = average_precision_score(y2_test, y3_test_proba)

print(f"\n'*80")
print("RESULTADOS DEL MODELO OPTIMIZADO")
print(f"\n'*80")

print(f"\n{'Metrica':<25} {'Train':<12} {'Test':<12} {'Diferencia':<12}")
print("-" * 65)
print(f"{'ROC-AUC':<25} {train3_auc:<12.4f} {test3_auc:<12.4f} {abs(train3_a}
print(f"\n{'Average Precision':<25} {train3_ap:<12.4f} {test3_ap:<12.4f} {abs(}

print(f"\nClassification Report (Test Set):")
print(classification_report(y2_test, y3_test_pred,
                           target_names=['Pago (0)', 'Default (1)']))

print(f"\nConfusion Matrix (Test Set):")
cm3_test = confusion_matrix(y2_test, y3_test_pred)
print(cm3_test)
print(f"\nInterpretacion:")
print(f" Verdaderos Negativos (TN): {cm3_test[0,0]},")
print(f" Falsos Positivos (FP): {cm3_test[0,1]},")
print(f" Falsos Negativos (FN): {cm3_test[1,0]},")
print(f" Verdaderos Positivos (TP): {cm3_test[1,1]},")
```

[4/5] Evaluando modelo con hiperparámetros óptimos...

```
=====
=====
=====
```

RESULTADOS DEL MODELO OPTIMIZADO

```
=====
=====
```

| Metrica | Train | Test | Diferencia |
|-------------------|--------|--------|------------|
| ROC-AUC | 0.9905 | 0.7482 | 0.2423 |
| Average Precision | 0.8949 | 0.2281 | 0.6668 |

Classification Report (Test Set):

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Pago (0) | 0.93 | 0.95 | 0.94 | 56538 |
| Default (1) | 0.29 | 0.24 | 0.26 | 4965 |
| accuracy | | | 0.89 | 61503 |
| macro avg | 0.61 | 0.59 | 0.60 | 61503 |
| weighted avg | 0.88 | 0.89 | 0.89 | 61503 |

Confusion Matrix (Test Set):

```
[[53668  2870]
 [ 3789 1176]]
```

Interpretacion:

Verdaderos Negativos (TN): 53,668
Falsos Positivos (FP): 2,870
Falsos Negativos (FN): 3,789
Verdaderos Positivos (TP): 1,176

```
In [44]: # =====
# COMPARACIÓN: MODELO 2 ORIGINAL VS MODELO OPTIMIZADO
# =====

# Extraer mejores hiperparámetros (si no están definidos)
best_params = random_search.best_params_

print("\n" + "="*80)
print("COMPARACIÓN: RF2 ORIGINAL vs RF3 OPTIMIZADO")
print("="*80)

print(f"\n{'Metrica':<25} {'RF2 Original':<15} {'RF3 Optimizado':<15} {'Mejor':<15}")
print("-" * 70)
print(f"{'ROC-AUC (Test)':<25} {test2_auc:<15.4f} {test3_auc:<15.4f} {test3_auc - test2_auc:<15.4f}")
print(f"{'Average Precision (Test)':<25} {test2_ap:<15.4f} {test3_ap:<15.4f} {test3_ap - test2_ap:<15.4f}")
print(f"{'ROC-AUC (Train)':<25} {train2_auc:<15.4f} {train3_auc:<15.4f} {train3_auc - train2_auc:<15.4f}")
print(f"{'Overfitting (Gap)':<25} {train2_auc - test2_auc:<15.4f} {train3_auc - test3_auc:<15.4f}")

print("\n" + "="*80)
print("HIPERPARÁMETROS: ORIGINAL vs OPTIMIZADO")
```

```

print("=*80)
print(f"\n{'Parámetro':<25} {'RF2 Original':<20} {'RF3 Optimizado':<20}")
print("-" * 65)
print(f"{'n_estimators':<25} {N_ESTIMATORS:<20} {best_params.get('n_estimators')}")
print(f"{'max_depth':<25} {MAX_DEPTH:<20} {str(best_params.get('max_depth'))}")
print(f"{'min_samples_split':<25} {MIN_SAMPLES_SPLIT:<20} {best_params.get('min_samples_split')}")
print(f"{'min_samples_leaf':<25} {MIN_SAMPLES_LEAF:<20} {best_params.get('min_samples_leaf')}")
print(f"{'max_features':<25} {'auto':<20} {str(best_params.get('max_features'))}")
print(f"{'criterion':<25} {'gini':<20} {best_params.get('criterion', 'gini')}")
print(f"{'bootstrap':<25} {'True':<20} {str(best_params.get('bootstrap', True))}")

# Análisis de mejora
diff_auc = test3_auc - test2_auc
diff_overfitting = (train3_auc - test3_auc) - (train2_auc - test2_auc)

print("\n" + "=*80)
print("ANÁLISIS DE MEJORA")
print("=*80)
if diff_auc > 0:
    print(f"\n✓ El modelo optimizado MEJORA el ROC-AUC en {diff_auc:.4f} ({diff_overfitting:.4f})")
else:
    print(f"\nx El modelo optimizado NO mejora el ROC-AUC ({diff_auc:.4f})")

if diff_overfitting < 0:
    print(f"\n✓ El overfitting se REDUCE en {abs(diff_overfitting):.4f}")
else:
    print(f"\nx El overfitting AUMENTA en {diff_overfitting:.4f}")

```

=====

====

COMPARACIÓN: RF2 ORIGINAL vs RF3 OPTIMIZADO

=====

====

| Metrica | RF2 Original | RF3 Optimizado | Mejora |
|--------------------------|--------------|----------------|---------|
| ROC-AUC (Test) | 0.7407 | 0.7482 | +0.0074 |
| Average Precision (Test) | 0.2133 | 0.2281 | +0.0148 |
| ROC-AUC (Train) | 0.9038 | 0.9905 | +0.0867 |
| Overfitting (Gap) | 0.1631 | 0.2423 | +0.0792 |

=====

====

HIPERPARÁMETROS: ORIGINAL vs OPTIMIZADO

=====

====

| Parámetro | RF2 Original | RF3 Optimizado |
|-------------------|--------------|----------------|
| n_estimators | 100 | 200 |
| max_depth | 15 | None |
| min_samples_split | 10 | 30 |
| min_samples_leaf | 5 | 5 |
| max_features | auto | sqrt |
| criterion | gini | entropy |
| bootstrap | True | True |

=====

====

ANÁLISIS DE MEJORA

=====

====

- ✓ El modelo optimizado MEJORA el ROC-AUC en 0.0074 (0.74%)
- ✗ El overfitting AUMENTA en 0.0792

In [46]: # =====

```
# =====
# VISUALIZACIONES DEL MODELO OPTIMIZADO
# =====
print("\n[5/5] Generando visualizaciones del modelo optimizado...")

# Extraer resultados de CV (si no están definidos)
cv_results = pd.DataFrame(random_search.cv_results_)
cv_results_sorted = cv_results.sort_values('rank_test_score').head(10)

# Calcular deciles para modelo optimizado
test_results3 = pd.DataFrame({
    'y_true': y2_test,
    'y_proba': y3_test_proba
})
test_results3['decil'] = pd.qcut(test_results3['y_proba'], q=10, labels=False

decil_stats3 = test_results3.groupby('decil').agg({
    'y_true': ['sum', 'count', 'mean']})
```

```

}).reset_index()
decil_stats3.columns = ['decil', 'positivos', 'total', 'tpr']
decil_stats3 = decil_stats3.sort_values('decil', ascending=False)

# Feature importance del modelo optimizado
feature_importance3 = pd.DataFrame({
    'Variable': X2.columns,
    'Importancia': best_model.feature_importances_
}).sort_values('Importancia', ascending=False)

# Crear figura comparativa
fig, axes = plt.subplots(2, 3, figsize=(20, 12))

# ROC Curves comparison
fpr2, tpr2, _ = roc_curve(y2_test, y2_test_proba)
fpr3, tpr3, _ = roc_curve(y2_test, y3_test_proba)

axes[0, 0].plot(fpr2, tpr2, linewidth=2, color='darkorange', label=f'RF2 Original')
axes[0, 0].plot(fpr3, tpr3, linewidth=2, color='green', label=f'RF3 Optimizada')
axes[0, 0].plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random')
axes[0, 0].set_xlabel('False Positive Rate', fontsize=12)
axes[0, 0].set_ylabel('True Positive Rate', fontsize=12)
axes[0, 0].set_title('ROC Curve Comparison', fontsize=14, fontweight='bold')
axes[0, 0].legend()
axes[0, 0].grid(alpha=0.3)

# Precision-Recall comparison
precision2, recall2, _ = precision_recall_curve(y2_test, y2_test_proba)
precision3, recall3, _ = precision_recall_curve(y2_test, y3_test_proba)

axes[0, 1].plot(recall2, precision2, linewidth=2, color='darkorange', label=f'RF2 Original')
axes[0, 1].plot(recall3, precision3, linewidth=2, color='green', label=f'RF3 Optimizada')
axes[0, 1].axhline(y=y2_test.mean(), color='gray', linestyle='--', label=f'Eje de la media')
axes[0, 1].set_xlabel('Recall', fontsize=12)
axes[0, 1].set_ylabel('Precision', fontsize=12)
axes[0, 1].set_title('Precision-Recall Comparison', fontsize=14, fontweight='bold')
axes[0, 1].legend()
axes[0, 1].grid(alpha=0.3)

# CV Results distribution
cv_scores_plot = cv_results_sorted['mean_test_score'].values[:10]
axes[0, 2].barh(range(10), cv_scores_plot, color='steelblue', alpha=0.7)
axes[0, 2].set_yticks(range(10))
axes[0, 2].set_yticklabels([f'Config {i+1}' for i in range(10)])
axes[0, 2].set_xlabel('ROC-AUC Score', fontsize=12)
axes[0, 2].set_title('Top 10 Configuraciones (CV)', fontsize=14, fontweight='bold')
axes[0, 2].invert_yaxis()
axes[0, 2].grid(alpha=0.3, axis='x')

# TPR por Decil comparison
width = 0.35
x = np.arange(len(decil_stats3))
axes[1, 0].bar(x - width/2, decil_stats2.sort_values('decil')['tpr'], width,
               axes[1, 0].bar(x + width/2, decil_stats3.sort_values('decil')['tpr'], width,
               axes[1, 0].set_xlabel('Decil de Score', fontsize=12)
               axes[1, 0].set_ylabel('TPR', fontsize=12)

```

```

axes[1, 0].set_title('TPR por Decil', fontsize=14, fontweight='bold')
axes[1, 0].set_xticks(x)
axes[1, 0].set_xticklabels(decil_stats3.sort_values('decil')['decil'].values)
axes[1, 0].legend()
axes[1, 0].grid(alpha=0.3, axis='y')

# Feature Importance (Top 15)
top_features3 = feature_importance3.head(15)
axes[1, 1].barh(range(len(top_features3)), top_features3['Importancia'].values)
axes[1, 1].set_yticks(range(len(top_features3)))
axes[1, 1].set_yticklabels(top_features3['Variable'].values)
axes[1, 1].set_xlabel('Importancia', fontsize=12)
axes[1, 1].set_title('Top 15 Variables (RF3 Optimizado)', fontsize=14, fontweight='bold')
axes[1, 1].invert_yaxis()
axes[1, 1].grid(alpha=0.3, axis='x')

# Comparison summary
summary_data = {
    'Modelo': ['RF2 Original', 'RF3 Optimizado'],
    'ROC-AUC Test': [test2_auc, test3_auc],
    'ROC-AUC Train': [train2_auc, train3_auc],
    'Overfitting': [train2_auc - test2_auc, train3_auc - test3_auc]
}
summary_df = pd.DataFrame(summary_data)

colors = ['darkorange', 'green']
x_pos = np.arange(len(summary_df))
width = 0.25

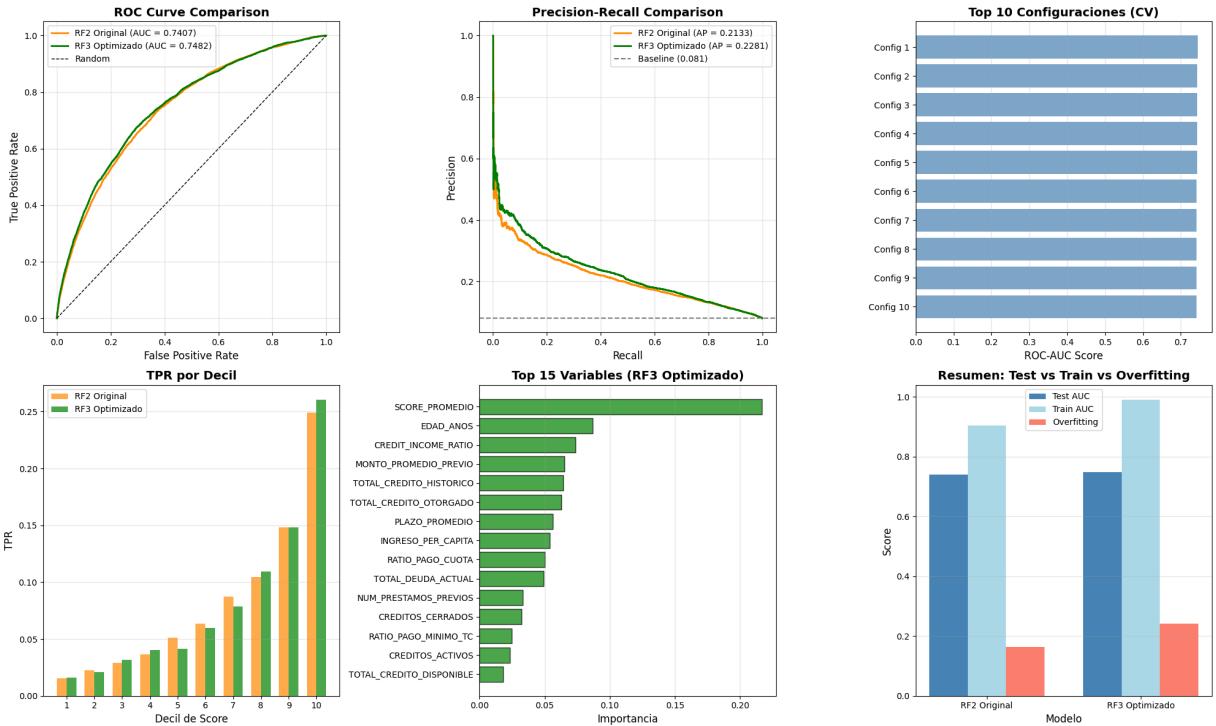
bars1 = axes[1, 2].bar(x_pos - width, summary_df['ROC-AUC Test'], width, label='Test', color=colors[0])
bars2 = axes[1, 2].bar(x_pos, summary_df['ROC-AUC Train'], width, label='Train', color=colors[1])
bars3 = axes[1, 2].bar(x_pos + width, summary_df['Overfitting'], width, label='Overfitting', color=colors[1])

axes[1, 2].set_xlabel('Modelo', fontsize=12)
axes[1, 2].set_ylabel('Score', fontsize=12)
axes[1, 2].set_title('Resumen: Test vs Train vs Overfitting', fontsize=14, fontweight='bold')
axes[1, 2].set_xticks(x_pos)
axes[1, 2].set_xticklabels(summary_df['Modelo'])
axes[1, 2].legend()
axes[1, 2].grid(alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig(f'{OUTPUT_DIR}/random_forest3_optimized_curves.png', dpi=300, bbox_inches='tight')
print(f"Gráficas guardadas: {OUTPUT_DIR}/random_forest3_optimized_curves.png")
plt.show()

```

[5/5] Generando visualizaciones del modelo optimizado...
 Gráficas guardadas: ../../data/random-forest-output/random_forest3_optimized_curves.png



```
In [ ]: # =====#
# GUARDAR MODELO OPTIMIZADO Y RESULTADOS
# =====#
print("\nGuardando modelo optimizado y resultados...")

# Guardar modelo optimizado
with open(f'{OUTPUT_DIR}/random_forest3_optimized_model.pkl', 'wb') as f:
    pickle.dump(best_model, f)

# Guardar mejores hiperparámetros
best_params_df = pd.DataFrame([best_params])
best_params_df.to_csv(f'{OUTPUT_DIR}/best_hyperparameters.csv', index=False)

# Guardar feature importance del modelo optimizado
feature_importance3.to_csv(f'{OUTPUT_DIR}/feature_importance_rf3_optimized.csv')

# Guardar predicciones del modelo optimizado
test_predictions3 = pd.DataFrame({
    'SK_ID_CURR': df.iloc[X2_test.index]['SK_ID_CURR'].values,
    'TARGET_Real': y2_test.values,
    'TARGET_Predicho': y3_test_pred,
    'Probabilidad_Default': y3_test_proba
})
test_predictions3.to_csv(f'{OUTPUT_DIR}/test_predictions_rf3_optimized.csv', index=False)

# Guardar resultados de CV
cv_results.to_csv(f'{OUTPUT_DIR}/cv_search_results.csv', index=False)

print(f"\nArchivos guardados en: {OUTPUT_DIR}/")
print(" - random_forest3_optimized_model.pkl")
print(" - best_hyperparameters.csv")
print(" - feature_importance_rf3_optimized.csv")
print(" - test_predictions_rf3_optimized.csv")
```

```
print(" - cv_search_results.csv")
print(" - random_forest3_optimized_curves.png")

# =====
# RESUMEN FINAL
# =====
print("\n" + "="*80)
print("RESUMEN FINAL – BÚSQUEDA DE HIPERPARÁMETROS COMPLETADA")
print("="*80)

print(f"\n📊 RESULTADOS DEL MODELO OPTIMIZADO (RF3)")
print(f"    ROC-AUC Test: {test3_auc:.4f}")
print(f"    Average Precision: {test3_ap:.4f}")
print(f"    Overfitting: {train3_auc - test3_auc:.4f}")

print(f"\n⭐ MEJORES HIPERPARÁMETROS:")
for param, value in best_params.items():
    print(f"    {param}: {value}")

print(f"\n〽️ COMPARACIÓN CON MODELO ORIGINAL:")
print(f"    Cambio en ROC-AUC: {test3_auc - test2_auc:+.4f}")
print(f"    Cambio en Overfitting: {(train3_auc - test3_auc) - (train2_auc - test2_auc)}")

print("\n" + "="*80)
print("¡OPTIMIZACIÓN DE HIPERPARÁMETROS COMPLETADA!")
print("="*80)
```